

Data visualization

R for Data Science

Juan R Gonzalez
`juanr.gonzalez@isglobal.org`

BRGE - Bioinformatics Research Group in Epidemiology
ISGlobal - Barcelona Institute for Global Health
<http://brge.isglobal.org>

Data distribution (categorical)


```
library(tidyverse)
diamonds
```

```
# A tibble: 53,940 x 10
  carat cut      color clarity depth table price      x      y      z
  <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>
1 0.23 Ideal    E      SI2     61.5   55   326   3.95   3.98   2.43
2 0.21 Premium  E      SI1     59.8   61   326   3.89   3.84   2.31
3 0.23 Good     E      VS1     56.9   65   327   4.05   4.07   2.31
4 0.290 Premium I      VS2     62.4   58   334   4.2    4.23   2.63
5 0.31 Good     J      SI2     63.3   58   335   4.34   4.35   2.75
6 0.24 Very Good J      VVS2     62.8   57   336   3.94   3.96   2.48
7 0.24 Very Good I      VVS1     62.3   57   336   3.95   3.98   2.47
8 0.26 Very Good H      SI1     61.9   55   337   4.07   4.11   2.53
9 0.22 Fair     E      VS2     65.1   61   337   3.87   3.78   2.49
10 0.23 Very Good H      VS1     59.4   61   338   4      4.05   2.39
# ... with 53,930 more rows
```

```
count(diamonds, cut)
```


```
# A tibble: 5 x 2
  cut      n
  <ord> <int>
1 Fair   1610
2 Good   4906
3 Very Good 12082
4 Premium 13791
5 Ideal  21551
```

```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut))
```



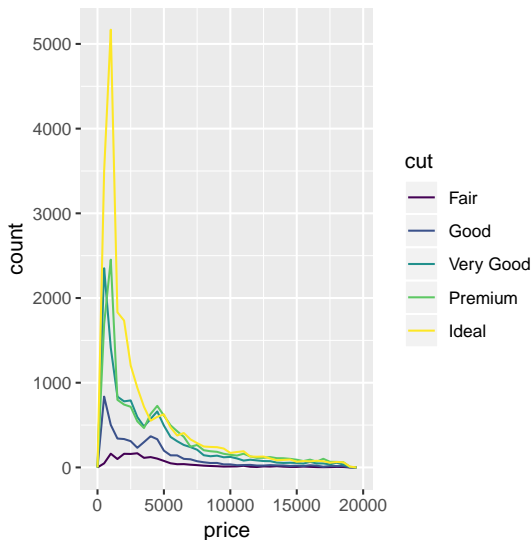
Data distribution (continuous)

```
ggplot(data = diamonds) +  
  geom_histogram(mapping = aes(x = carat), binwidth = 0.5)
```



Data distribution (categorical and continuous)

```
ggplot(data = diamonds, mapping = aes(x = price)) +  
  geom_freqpoly(mapping = aes(colour = cut), binwidth = 500)
```



Data distribution (categorical and continuous)

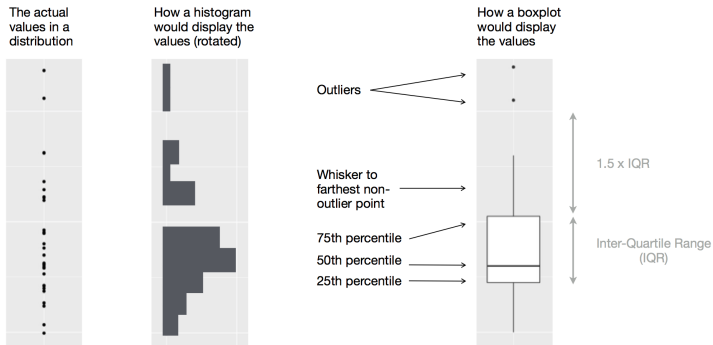



Figure 1: Box-plot


Boxplot

```
ggplot(data = mpg, mapping = aes(x = class, y = hwy)) +  
  geom_boxplot()
```




Reordered boxplot (improves interpretation)

```
ggplot(data = mpg) +  
  geom_boxplot(mapping = aes(x = reorder(class, hwy,  
                                FUN = median),  
                              y = hwy))
```





```
ggplot(data = mpg) +  
  geom_boxplot(mapping = aes(x = reorder(class, hwy,  
                                     FUN = median),  
                             y = hwy)) + coord_flip()
```




Improved boxplot

```
ggplot(iris, aes(x=Species, y=Sepal.Width) ) +  
  geom_boxplot(alpha=0.3, outlier.colour = "blue") +  
  geom_point(stat= "summary", fun.y=mean,  
            shape=16, size=1.5, color="red") +  
  geom_jitter(width = 0.1, alpha = 0.2)
```



Two categorical variables

```
ggplot(data = diamonds) +  
  geom_count(mapping = aes(x = cut, y = color))
```




Another approach is to compute the count with dplyr:

```
diamonds %>%  
  count(color, cut)
```

```
# A tibble: 35 x 3  
  color cut      n  
  <ord> <ord> <int>  
1 D    Fair    163  
2 D    Good    662  
3 D    Very Good 1513  
4 D    Premium 1603  
5 D    Ideal   2834  
6 E    Fair    224  
7 E    Good    933  
8 E    Very Good 2400  
9 E    Premium 2337  
10 E   Ideal   3903  
# ... with 25 more rows
```


and then visualize with `geom_tile()` and the fill aesthetic:

```
diamonds %>%  
  count(color, cut) %>%  
  ggplot(mapping = aes(x = color, y = cut)) +  
    geom_tile(mapping = aes(fill = n))
```




Two continuos variables

```
ggplot(data = diamonds) +  
  geom_point(mapping = aes(x = carat, y = price))
```




Scatterplots become less useful as the size of your dataset grows because points begin to overplot: use alpha aesthetic:

```
ggplot(data = diamonds) +  
  geom_point(mapping = aes(x = carat, y = price),  
              alpha = 1 / 100)
```




Another option is to bin one continuous variable and use boxplot

```
ggplot(data = diamonds, mapping = aes(x = carat, y = price)) +  
  geom_boxplot(mapping = aes(group = cut_width(carat, 0.1)))
```



Fitting regression line

```
library(ggpmisc)
set.seed(1234)
iris <- mutate(iris,
               Y = 1.5 + 3.2*Sepal.Width +
                 rnorm(nrow(iris)))
ggplot(iris, aes(x = Sepal.Width, y = Y)) +
  geom_smooth(method = "lm", se=FALSE, color="black",
             formula = y ~ x) +
  stat_poly_eq(formula = y ~ x,
              aes(label = paste(..eq.label.., ..rr.label..,
                                sep = "~~~")),
              parse = TRUE) +
  geom_point()
```



Why using ggplot2?


Let's imagine we want to compare gene expression by growth rate in twenty genes in six conditions¹ example obtained from <http://varianceexplained.org/r/why-i-use-ggplot2/>)

```
load("../..data/genes.Rdata")
genes
```

```
# A tibble: 711 x 7
  name BP          MF          systematic_name nutrient rate expression
  <chr> <chr>        <chr>        <chr>        <chr>   <dbl>   <dbl>
1 SUL1 sulfate tran~ sulfate trans~ YBR294W      Glucose  0.05   -0.32
2 ""    biological p~ molecular fun~ YKL187C      Glucose  0.05    4.13
3 QDR2 multidrug tr~ multidrug eff~ YIL121W      Glucose  0.05    1.07
4 LEU1 leucine bios~ 3-isopropylma~ YGL009C      Glucose  0.05   -1.12
5 PHO5 phosphate me~ acid phosphat~ YBR093C      Glucose  0.05    2.39
6 PHO12 biological p~ acid phosphat~ YHR215W      Glucose  0.05    0.9
7 PHO11 phosphate me~ acid phosphat~ YAR071W      Glucose  0.05    1.14
8 GIT1 glycerophosph~ glycerophosph~ YCR098C      Glucose  0.05    0.77
9 AGP3 amino acid t~ amino acid tr~ YFL055W      Glucose  0.05    0.570
10 ""    biological p~ molecular fun~ YOL164W      Glucose  0.05    0.53
# ... with 701 more rows
```

¹(

```
ggplot(genes, aes(rate, expression, color = nutrient)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  facet_wrap(~name + systematic_name, scales = "free_y")
```



```

par(mar = c(1.5, 1.5, 1.5, 1.5))

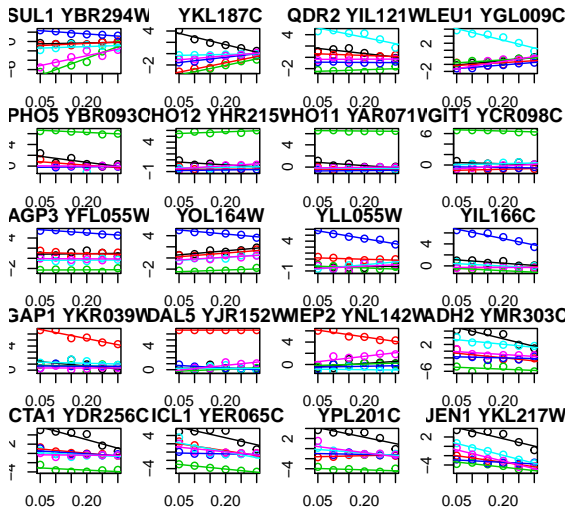
colors <- 1:6
names(colors) <- unique(genes$nutrient)

m <- matrix(c(1:20, 21, 21, 21, 21), nrow = 6,
             ncol = 4, byrow = TRUE)
layout(mat = m, heights = c(.18, .18, .18, .18, .18, .1))

genes$combined <- paste(genes$name, genes$systematic_name)
for (gene in unique(genes$combined)) {
  sub_data <- filter(genes, combined == gene)
  plot(expression ~ rate, sub_data,
        col = colors[sub_data$nutrient], main = gene)
  for (n in unique(sub_data$nutrient)) {
    m <- lm(expression ~ rate,
             filter(sub_data, nutrient == n))
    if (!is.na(m$coefficients[2])) {
      abline(m, col = colors[n])
    }
  }
}

# create a new plot for legend
plot(1, type = "n", axes = FALSE, xlab = "", ylab = "")
legend("top", names(colors), col = colors, horiz = TRUE, lwd = 4

```




ggplot2

Every visualization in ggplot2 is composed of the following:


- ▶ Data: the raw material of your visualization
- ▶ Layers: what you see on the plot (e.g., points, lines, ...)
- ▶ Scales: Maps the data to graphical output
- ▶ Coordinates: The visualization's perspective (e.g., grids)
- ▶ Faceting: Provides “visual drill-down” into the data
- ▶ Themes: Controls the details of the display (e.g. fonts)

plot = data + Aesthetics + geometry

ggplot2 is based on the **grammar of graphics**, the idea that you can build every graph from the same few components: a **data** set, a set of **geoms**—visual marks that represent data points, and a **coordinate system**.



To display data values, map variables in the data set to aesthetic properties of the geom like **size**, **color**, and **x** and **y** locations.



Some plots visualize a **transformation** of the original data set. Use a **stat** to choose a common transformation to visualize, e.g. `a + geom_bar(stat = "bin")`





Figure 2: ggplot cheat sheet:

<https://www.rstudio.com/resources/cheatsheets/>


```
mtcars[1:5,1:8]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0

```
ggplot(mtcars) +           # data  
  aes(x = mpg, y=wt) +     # Aesthetics  
  geom_point()             # geometry (layer)
```




Possible aesthetics

These are some:

- ▶ `theme_dark()`
- ▶ `theme_minimal()`
- ▶ `theme_classic()`
- ▶ `theme_void()`
- ▶ `theme_test()`

```
ggplot(mtcars) +           # data
  aes(x = mpg, y=wt) +     # Aesthetics
  geom_point() +           # geometry (layer)
  theme_minimal()          # theme
```



Possible geometry (layers)

- ▶ For one continuous variable:
 - ▶ `geom_area()` for area plot
 - ▶ `geom_density()` for density plot
 - ▶ `geom_dotplot()` for dot plot
 - ▶ `geom_freqpoly()` for frequency polygon
 - ▶ `geom_histogram()` for histogram plot
 - ▶ `stat_ecdf()` for empirical cumulative density function
 - ▶ `stat_qq()` for quantile - quantile plot
- ▶ For one discrete variable:
 - ▶ `geom_bar()` for bar plot

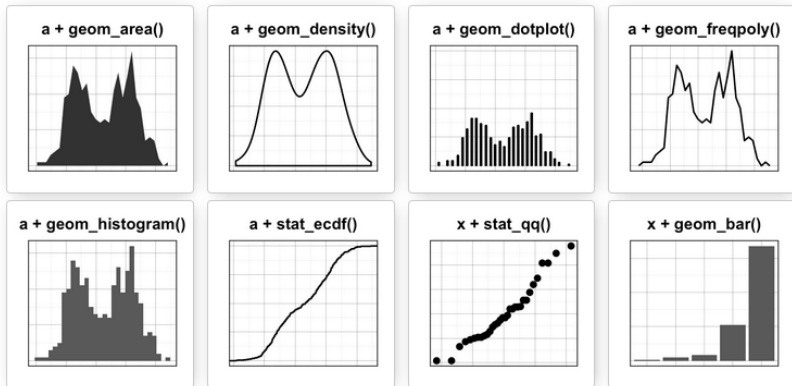



Figure 3: Possible layers

```
ggplot(mtcars) +  
  aes(sample = mpg) + stat_qq()
```




Facets

```
data(tips, package="reshape2")  
head(tips)
```


	total_bill	tip	sex	smoker	day	time	size
1	16.99	1.01	Female	No	Sun	Dinner	2
2	10.34	1.66	Male	No	Sun	Dinner	3
3	21.01	3.50	Male	No	Sun	Dinner	3
4	23.68	3.31	Male	No	Sun	Dinner	2
5	24.59	3.61	Female	No	Sun	Dinner	4
6	25.29	4.71	Male	No	Sun	Dinner	4


```
sp <- ggplot(tips, aes(x=total_bill, y=tip/total_bill)) +  
  geom_point()  
sp
```




Facet grid


```
# vertical direction  
sp + facet_grid(sex ~ .)
```



```
# horizontal direction  
sp + facet_grid(. ~ sex)
```



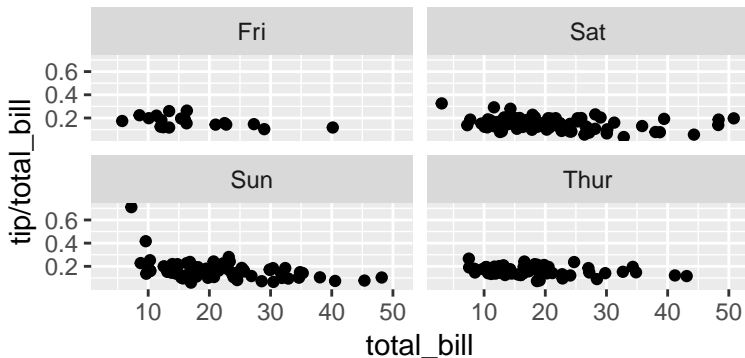
```
# Divide with "sex" vertical, "day" horizontal  
sp + facet_grid(sex ~ day)
```



Facet wrap

Instead of faceting with a variable in the horizontal or vertical direction, facets can be placed next to each other, wrapping with a certain number of columns or rows. The label for each plot will be at the top of the plot.


```
# Divide by day, going horizontally and wrapping with 2 columns  
sp + facet_wrap( ~ day, ncol=2)
```




labeller() can use any function that takes a character vector as input and returns a character vector as output (e.g. gsub). We can also define our own custom functions, like this one, which reverses strings:

```
# Reverse each strings in a character vector
reverse <- function(strings) {
  strings <- strsplit(strings, "")
  vapply(strings, function(x) {
    paste(rev(x), collapse = "")
  }, FUN.VALUE = character(1))
}

sp + facet_grid(. ~ sex, labeller=labeller(sex = reverse))
```



```
sp + facet_grid(sex ~ day) +
  theme(strip.text.x = element_text(size=8, angle=75),
        strip.text.y = element_text(size=12, face="bold"),
        strip.background = element_rect(colour="brown",
                                         fill="tomato"))
```



More about ggplot2

<https://ggplot2.tidyverse.org/reference/>

Exercises (data visualization)

1. Use what you've learned to improve the visualisation of the departure times of cancelled vs. non-cancelled flights (NOTE: missing values in the `dep_time` variable indicate that the flight was cancelled).
2. What variable in the diamonds dataset is most important for predicting the price of a diamond? How is that variable correlated with cut? Why does the combination of those two relationships lead to lower quality diamonds being more expensive?. NOTE: variable cut describes the quality (see `?diamonds`).
3. Visualize the number of flights of each airline by month.
4. Instead of summarising the conditional distribution with a boxplot, you could use a frequency polygon. What do you need to consider when using `cut_width()` vs `cut_number()`? How does that impact a visualisation of the 2d distribution of carat and price (diamonds dataset)?
5. Visualise the distribution of carat, partitioned by price on diamonds dataset .

6. Download genome data (https://github.com/isglobal-brge/TeachingMaterials/blob/master/Master_Modelling/data/genome.txt) into your computer and load it into RStudio by using `read_delim` function (NOTE: data are tab-delimited).

The file is containing information about

- ▶ Name: genomic variant (single nucleotide polymorphism)
- ▶ Chr: chromosome
- ▶ Position: chromosome position
- ▶ Log.R.Ratio: log-ratio intensity of the two alleles
- ▶ B.Allele.Freq: frequency of the alternative allele
 - ▶ 6.1 Which is the expected value of Log.R.Ratio and B.Allele.Freq for each chromosome? (show the R code to get your answer)
 - ▶ 6.2 Create a facet plot that represent the Log.R.Ratio for each chromosome
 - ▶ 6.3 Create a facet plot of B.Allele.Freq for chromosomes 1, 2, 3, ..., 6 drawing B.Allele.Freq information in red color.

Session info

sessionInfo()

R version 3.5.0 (2018-04-23)

Platform: x86_64-w64-mingw32/x64 (64-bit)

Running under: Windows 10 x64 (build 17134)

Matrix products: default

locale:

[1] LC_COLLATE=Spanish_Spain.1252 LC_CTYPE=Spanish_Spain.1252

[3] LC_MONETARY=Spanish_Spain.1252 LC_NUMERIC=C

[5] LC_TIME=Spanish_Spain.1252

attached base packages:

[1] stats graphics grDevices utils datasets methods base

other attached packages:

[1] bindrcpp_0.2.2 ggpmisc_0.3.0 forcats_0.3.0 stringr_1.3.1

[5] dplyr_0.7.6 purrr_0.2.4 readr_1.1.1 tidyr_0.8.0

[9] tibble_1.4.2 ggplot2_3.0.0 tidyverse_1.2.1

loaded via a namespace (and not attached):

[1] tidyselect_0.2.4 reshape2_1.4.3 haven_1.1.1 lattice_0.20-35

[5] colorspace_1.3-2 viridisLite_0.3.0 htmltools_0.3.6 yaml_2.1.19

[9] utf8_1.1.3 rlang_0.2.2 pillar_1.2.2 foreign_0.8-70

[13] glue_1.2.0 withr_2.1.2 modelr_0.1.2 readxl_1.1.0

[17] bindr_0.1.1 plyr_1.8.4 munsell_0.5.0 gtable_0.2.0

[21] cellranger_1.1.0 rvest_0.3.2 codetools_0.2-15 psych_1.8.4

[25] evaluate_0.10.1 labeling_0.3 knitr_1.20 parallel_3.5.0

[29] broom_0.4.4 Rcpp_0.12.18 polynom_1.3-9 scales_1.0.0

[33] backports_1.1.2 jsonlite_1.5 mnormt_1.5-5 hms_0.4.2